

Assignment 3: Meeting Transcript Summarizer & Action-Item Extractor

1. Background & Context

Teams often record meetings and distribute raw transcripts, but reading through hundreds of lines to find key takeaways and next steps is time-consuming. Automating this process with a Generative AI model can save hours and ensure no action item is missed.

2. Problem Statement

“Given a plain-text meeting transcript, automatically produce (a) a concise two-sentence summary of the meeting’s main discussion points, and (b) a list of all action items mentioned, each prefaced with a dash (‘-’), outputting the result as a JSON object.”

3. Objectives & Learning Outcomes

- Prompt Design: Craft clear instructions to elicit structured responses (summary + action list) from an LLM.
- Structured Output: Enforce a JSON schema so downstream tools can reliably ingest the result.
- Error Handling: Detect and recover from malformed LLM outputs (e.g. missing JSON).
- Minimal UI/CLI: (Optional) Wrap the extractor in a simple script or notebook to demo end-to-end.

4. Technical Requirements

1. Input:

- A single string containing the full meeting transcript (100–600 words).
- Speaker labels (e.g. “Alice: ...”, “Bob: ...”) may be present but are optional.

2. Output:

- A JSON object with exactly two keys:

```
{  
  "summary": "...",  
  "action_items": [  
    "- ...",  
    "- ..."  
  ]  
}
```

3. Model:

- Use GPT-4o-mini (or GPT-4.1-mini) with temperature=0 for consistency.
- Wrap your prompt in a function, e.g.
def extract_meeting_notes(transcript: str) -> dict:
 # call OpenAI API with prompt template...
 return parsed_json

4. Validation:

- Define a Pydantic model to enforce:
class MeetingNotes(BaseModel):
 summary: str
 action_items: List[str]
- On parser errors, retry once with a “Please output valid JSON only” system prompt.

5. Prompt Template (Example)

You are a meeting assistant.

1. Summarize the meeting transcript below in exactly two sentences.
2. Then list all action items mentioned, each as a separate bullet beginning with a dash.

Return the result strictly as JSON with keys "summary" and "action_items".

Transcript:

```
{transcript}
```

6. Deliverables

1. Code: Jupyter Notebook OR zip of .py files.

7. Sample Inputs for Testing

Sample Input 1

Alice: Welcome everyone. Today we need to finalize the Q3 roadmap.

Bob: I've emailed the updated feature list—please review by Friday.

Carol: I'll set up the user-testing sessions next week.

Dan: Let's push the new UI mockups to staging on Wednesday.

Alice: Great. Also, can someone compile the stakeholder feedback into a slide deck?

Bob: I can handle the slide deck by Monday.

Alice: Thanks, team. Meeting adjourned.

Sample Input 2

Host: Let's kick off our marketing sync.

Emma: The social campaign draft is 80% done; I'll share it today.

Frank: I spoke with the design team—they'll deliver assets by Tuesday.

Emma: Once we have assets, I'll schedule the ads for next week.

George: Reminder: submit your budget requests before end of day.

Host: Noted. I'll send out the final budget spreadsheet.

Sample Input 3

John: Sales numbers for April are in-up 12% over target.

Lia: I'll distribute the detailed report to leadership this afternoon.

Mike: We need to train two new reps to handle increased volume.

Lia: I'll coordinate recruiting with HR by end of week.

John: Excellent. Let's aim to onboard them by May 15th.

Sample Input 4

Sample Input 5

HR: We need to update the travel-expense policy.

Alex: I'll draft the new policy doc and share with legal.

HR: Please include guidelines on per diem limits.

Legal: I'll review and provide feedback within three days.

HR: Once finalized, schedule a companywide announcement.